

Implementing Modern Fast Multipole Methods

Srinath Kailasa

Department of Mathematics
University College London

March 5, 2024



Table of Contents

Problem Setting

High Performance Computing

Concluding Remarks

Table of Contents

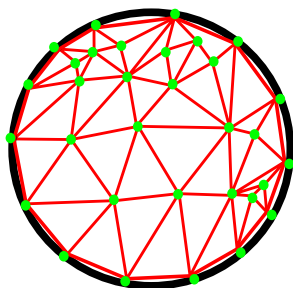
Problem Setting

High Performance Computing

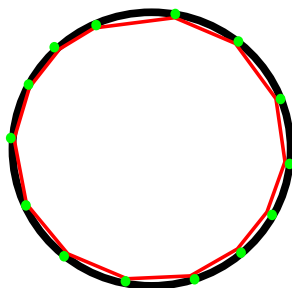
Concluding Remarks

Numerical Methods for PDEs

Consider 2D problem:



FEM



BEM

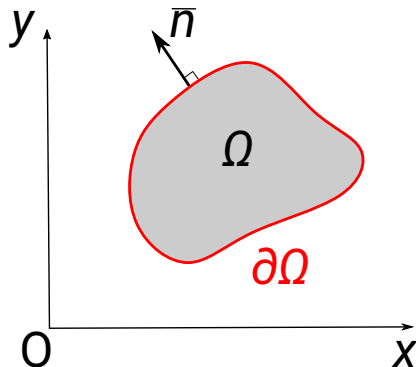
Laplace Interior Boundary Value Problem I

Want to solve

$$\Delta u(x) = 0 \quad (1)$$

For $x \in \Omega$, with boundary $\partial\Omega$.

Form an *integral equation* using
representation formula +
boundary data.



We seek $u \in C^2(\Omega) \cap C^2(\bar{\Omega})$ which satisfies $u = f$ on $\partial\Omega$.
Let's write down a solution using a 'double layer potential',

$$u(x) = K\phi = \int_{\partial\Omega} \phi(y) \frac{\partial\Phi(x,y)}{\partial n(y)} ds(y), \quad x \in \Omega$$

$\phi(y)$ unknown surface density, $\Phi(x,y)$ fundamental solution. End up with *boundary* integral equation,

$$\phi(x) - 2 \int_{\Gamma} \phi(y) \frac{\partial\Phi(x,y)}{\partial n(y)} ds(y) = -2f(x), \quad x \in \Gamma$$

$$-\frac{1}{2}\phi + K\phi = f$$

of the form

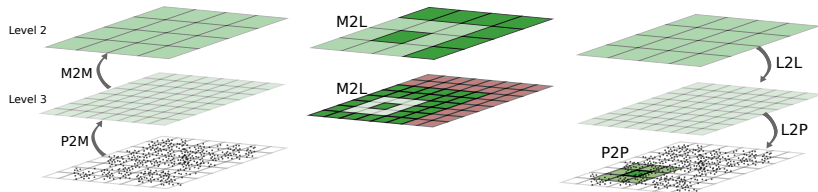
$$A\phi = f$$

with $A = K - \frac{1}{2}I$

Problem: A is *dense*. Slow $O(N^2)$ application cost, $O(N^3)$ inversion cost. CPUs of the 1980s can't handle this.
... But get to automatically incorporate boundary data *and* get a reduction in problem dimension.

Solution: Hierarchical Solvers (Fast Multipole Method). Can apply in $O(N)$ (sometimes), and invert using Krylov methods. Workaround for slow computers!

Fast Multipole Methods I



Consider N source and target particles,

- $O(N)$ boxes, so at least $O(N \log N)$ complexity.

FMM for Laplace is just $O(N)$, as number of boxes in interaction lists limited to constant number.

Fast Multipole Methods II

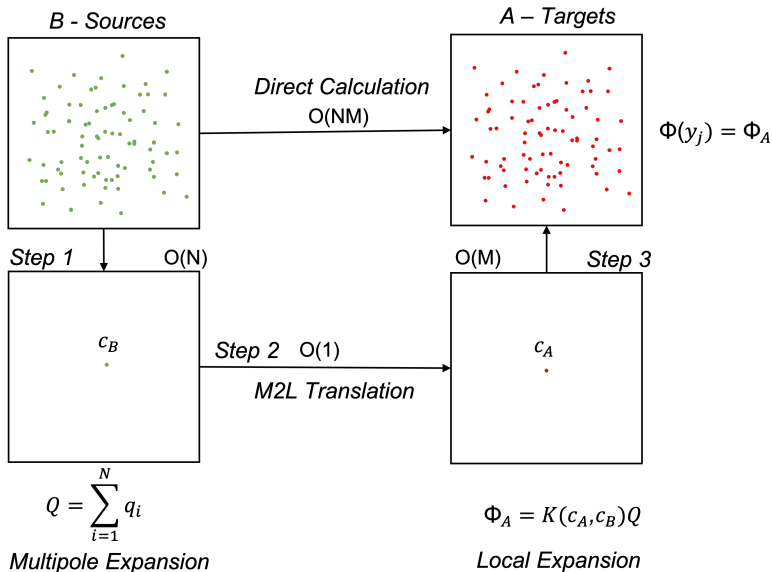


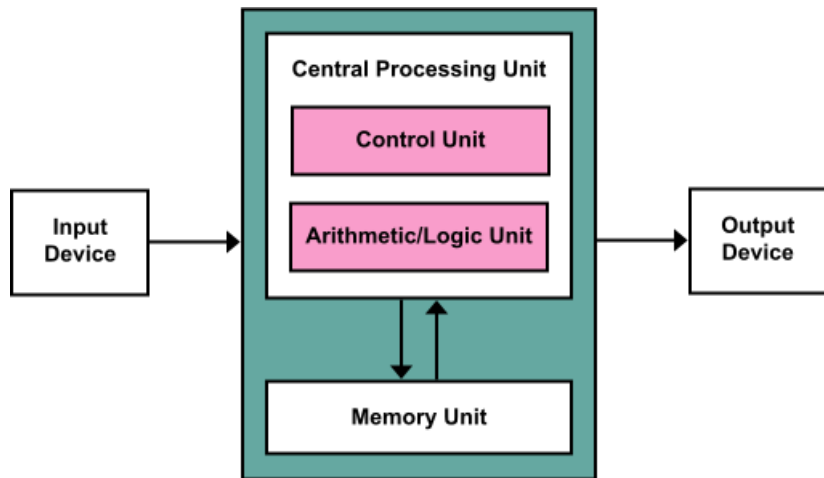
Table of Contents

Problem Setting

High Performance Computing

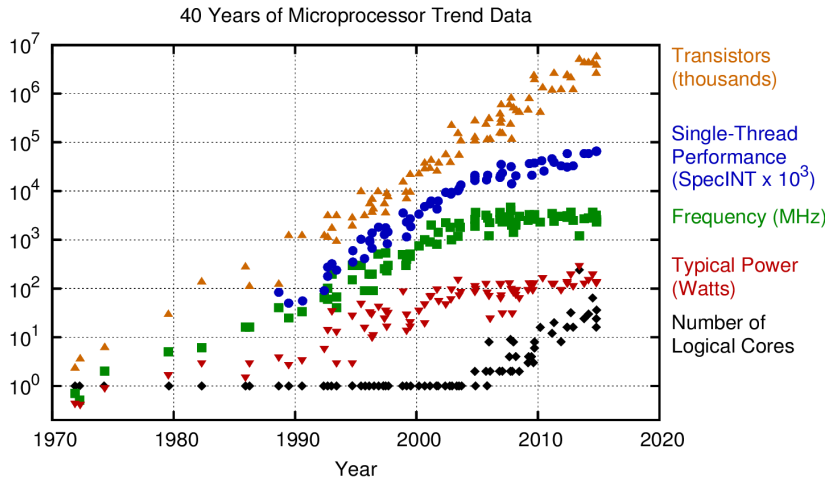
Concluding Remarks

Computer History I



source: [Wikipedia: Von Neumann Architectures](#)

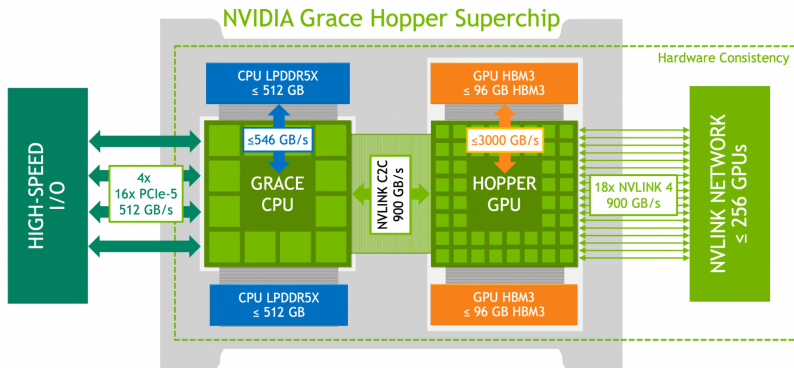
Computer History II



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

source: Conte, T, IEEE (2015)

Computer History II

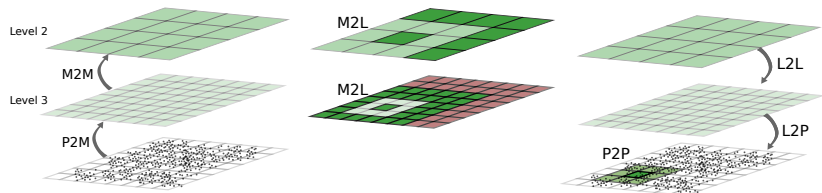


source: NVidia Corporation

Computational Bottlenecks for the FMM

- **P2P** - Direct Kernel evaluations - $O(N^2)$ flops $O(N)$ memory accesses.
- **M2L** - Translation of the Multipole into a Local Expansion (Formulation Dependant).

Computational Bottlenecks for the FMM II



The Kernel Independent FMM I

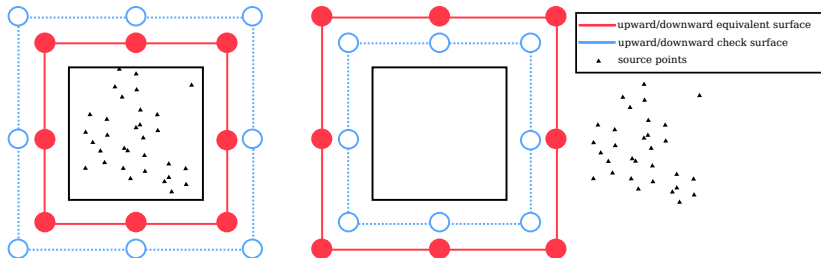
Consider simple test case, want to evaluate Laplace Green's function between N targets $\{x_i\}_{i=1}^N$ and N sources $\{y_i\}_{i=1}^N$.

$$\phi(x_i) = \sum_{j=1}^N K(x_i, y_j) q(y_j) \quad (2)$$

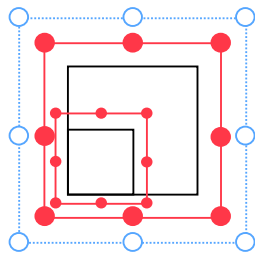
$$K(x, y) = \frac{1}{4\pi \|x - y\|}$$

Want to accelerate with an FMM, to compute in $O(N)$

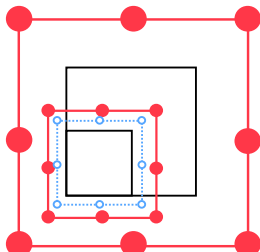
The Kernel Independent FMM II



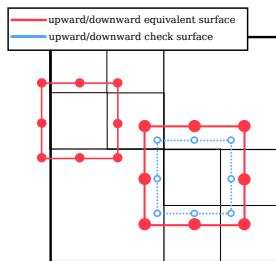
The Kernel Independent FMM III



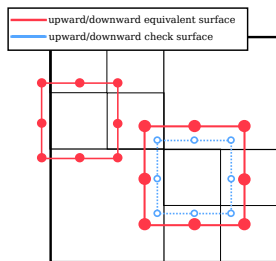
M2M



L2L



M2L



Can view M2L as either a convolution (FFT) or as a dense matrix vector product (BLAS)

$$\phi^{B,d}(x) = (K * q)(x) = \int_{y^{B,d}} K(x - y) q^{A,u}(y) dy \quad (3)$$

FFT memory throttled (notice the element-wise product), but is of only $O(N \log N)$ flops.

$$\phi(x) = \mathcal{F}^{-1}\{\mathcal{F}\{K\}(k) \cdot \mathcal{F}\{q\}(k)\}(x) \quad (4)$$

BLAS is compute bound, but of $O(p^2)$ flops ...

Modern Implementations I

Modern implementations of algorithms have to reflect on advances in hardware, where we are memory bound and flops are *cheap*.

1 **DESIGNING KERNEL INDEPENDENT FAST MULTIPOLE**
2 **METHODS FOR MODERN ARCHITECTURES ***

3 SRINATH KAILASA[†] AND TIMO BETCKE[‡]

4 **Abstract.** Kernel independent Fast Multipole Methods (kiFMM), with their broad applicability
5 to a wide variety of elliptic partial differential equations with asymptotically smooth kernels, have
6 emerged as a preferred method for computer implementations of recent benchmark softwares due to

Modern Implementations II

- Counterintuitively prefer a dense matrix vector product to an FFT, because of greater arithmetic intensity.
- Can stack together vectors which share a matrix, i.e. matrix multiplication, and get even greater data re-use.
- Can use numerical compression (SVD etc) to reduce size of (low-rank) matrices.
- Use specialised hardware and unified memory to keep the cost of GPU-CPU data exchange minimal.

Table of Contents

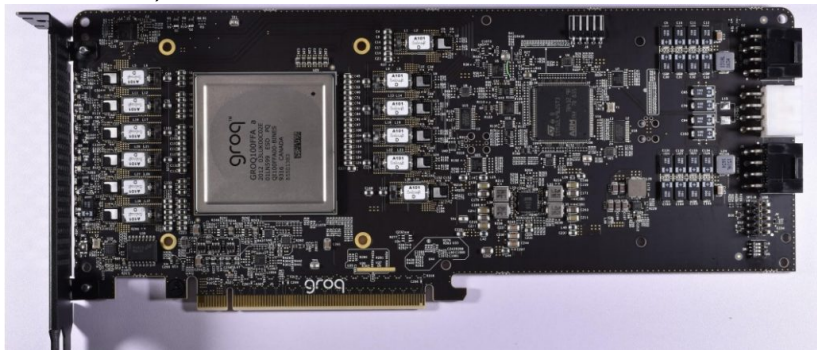
Problem Setting

High Performance Computing

Concluding Remarks

Conclusion

Key Takeaway: Hardware, software and algorithms advance in lockstep. Who knows what the computing systems of the future will enable! (Or the algorithms that will be devised to circumvent their failings).



source: [The Next Platform \(2023\)](#)