

Implementing Fast Multipole Methods with High Level Interpreted Languages

Srinath Kailasa

Department of Physics & Astronomy
University College London

September 9, 2020

Table of Contents

Fast Multipole Methods (FMMs)

PyExaFMM

Research Context

Table of Contents

Fast Multipole Methods (FMMs)

- Motivation

- Analytic FMM

- Kernel Independent FMM

PyExaFMM

Research Context

Motivation - the N Body Problem

- e.g. Electrostatics, Gravitation
- $\{x_i\}_{i=1,\dots,N}$ Source Particles
- $\{y_j\}_{j=1,\dots,M}$ Target Particles

$$\Phi(y_j) = \sum_{i=1}^N K(x_i, y_j) q_i, \quad (1)$$

$$\text{where, } K(x, y) = \frac{1}{\|x - y\|} \quad (2)$$

- FMM reduces complexity from $O(N^2)$ to $O(N)$



Motivation - Intuition




Galaxy A



Galaxy B

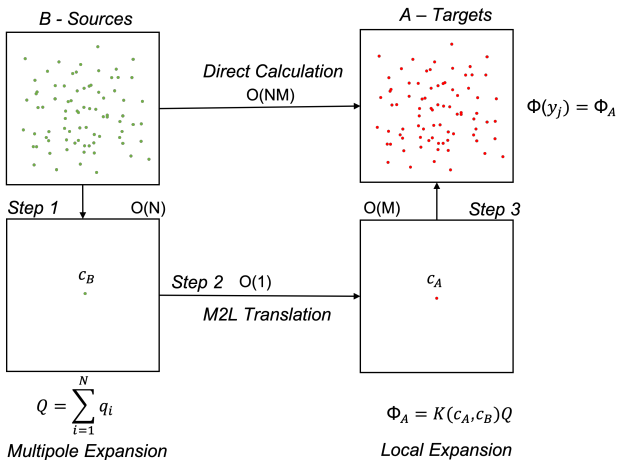


Motivation - Intuition



UCL


Motivation - Three Step Procedure




Analytic FMM - Concept

Idea: Use compressed representations of far field potentials to reduce complexity, in a recursive fashion

Analytic FMM - Multipole Expansion



Analytic FMM - Local Expansion




- For 3D Laplace kernel, can write multipole and local expansions using sph. harmonics, these can be truncated to required accuracy
- Exact operators exist for this kernel to shift between multipole and local expansion coefficients
- Exact bounds on error also exist, with respect to direct computation [2]


Analytic FMM - Motivating Problem

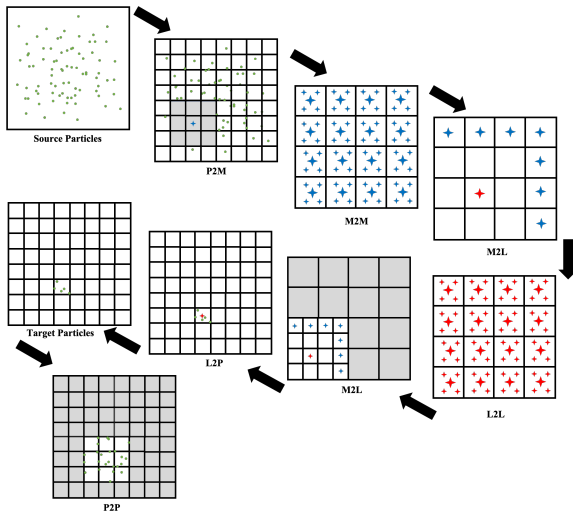
- Consider 2D Problem
- Domain, $\Omega = [0, 1] \times [0, 1]$
- Partition into recursively defined Quadtree
- Each level, l , partitioned into 4^l boxes
- Source and Target particles taken to be the same

Analytic FMM - Shifting Multipole Expansion



Analytic FMM - Shifting Local Expansion






Analytic FMM - Implementation Issues

- Representing problem with efficient data structures:
Quad/Octrees
- Computing and storing new expansions for each kernel, may require new software implementations

Kernel Independent FMM


- KIFMM only requires kernel evaluations
- Works by matching potential generated by particles with that generated by an equivalent density in the far field

Kernel Independent FMM - Least Squares Problem



Adapted from [3]

Kernel Independent FMM - Least Squares Problem



Adapted from [3]

Kernel Independent FMM - Least Squares Problem

- Check Surface $x^{B,u}$
- Equivalent Surface $y^{B,u}$
- Equivalent Density $\phi^{B,u}$
- Check Potential $q^{B,u}$
- Indices of source points I_s^B
- Source densities ϕ_i

$$\int_{y^{B,u}} K(x, y) \phi^{B,u} dy = \sum_{i \in I_s^B} K(x, y) \phi_i = q^{B,u} \text{ for any } x \in x^{B,u} \quad (3)$$




UCL

Kernel Independent FMM - Least Squares Problem

$$K_A \phi^A = K_B \phi^B \quad (4)$$

$$\phi^A = (\alpha I + K_A^* K_A)^{-1} K_B \phi^B \quad (5)$$

Kernel Independent FMM - Least Squares Problem, M2L



UCL

Kernel Independent FMM - Least Squares Problem, M2L

- Check Surface $x^{B,u}$
- Downward Equivalent Surface $y^{B,d}$
- Upward Equivalent Surface $y^{A,u}$
- Downward Equivalent Density $\phi^{B,d}$
- Upward Equivalent Density $\phi^{A,u}$

$$\int_{y^{A,u}} K(x, y) \phi^{A,u} dy = \int_{y^{B,d}} K(x, y) \phi^{B,d} dy, \text{ for any } x \in x^{B,d} \quad (6)$$

Kernel Independent FMM - Implementation Issues

- No need for new software implementation for large class of compatible Kernels
- Can built singly, extensible, and optimisable software implementation

Table of Contents

Fast Multipole Methods (FMMs)

PyExaFMM

Motivation

Goals

Outcomes

Research Context

Motivation

- Python has emerged as a standard in scientific and data intensive computing
- Desire a high quality software implementation which is also highly performant and easily portable
- Tradeoff performance of compiled languages for engineering ease, and portability

Goals

- Create a performant 3D Python implementation of the KIFMM
- Write software in an extensible and well tested way
- Take advantage of distributed and parallel computing concepts as much as possible


Outcomes - Vectorised Data Structures

10	11
00	01


Outcomes - Vectorised Data Structures

			11 11
00 00			

Outcomes - Vectorised Data Structures



Outcomes - Vectorised Data Structures



Outcomes - JIT Compilation

- Pre-analyse interpreted bytecode for reused calculations, compile into assembly and cache
- Use Numba on top of Numpy Containers

Outcomes - Multiprocessing & Operator caching

- Pre-compute and cache all operators, used in rhs of least squares problem, $(\phi^A = (\alpha I + K_A^* K_A)^{-1} K_B \phi^B)$
- Use process level parallelism to distribute the computation
- Use HDF5 for rapid loading to memory in comparison to simple serialisation

Outcomes - Low-Rank SVD Compression

- I - number of source boxes in interaction list of a given target box

$$K_{\text{source}} = (\alpha I + K_{\text{source}}^* K_{\text{source}})^{-1} K_{\text{target}} \quad (7)$$

$$K_{\text{concatenated}} = [K_1 | K_2 | \dots | K_I] \quad , \text{ where, } \{K_i | i \in [1, 2, \dots, I]\} \quad (8)$$

Outcomes - Low-Rank SVD Compression

- The rank of the full concatenated matrix is r
- Can truncate sum to target rank, k , and tune to retain most of the action of the matrix

$$K_{\text{concatenated}} = U\Sigma V^* \quad (9)$$

$$K_{\text{concatenated}} = \sum_{j=1}^r \sigma_j u_j v_j^* \quad (10)$$

$$\hat{K}_{\text{concatenated}} = \sum_{j=1}^k \sigma_j u_j v_j^* \quad (11)$$


Outcomes - Extensible Software Design

- Full unit tests
- Implements dependency inversion for operator loading
- Implements separation of concerns, between optimisation and logic



UCL

Outcomes - Complexity Achieved



Outcomes - Optimum Target Rank Investigated

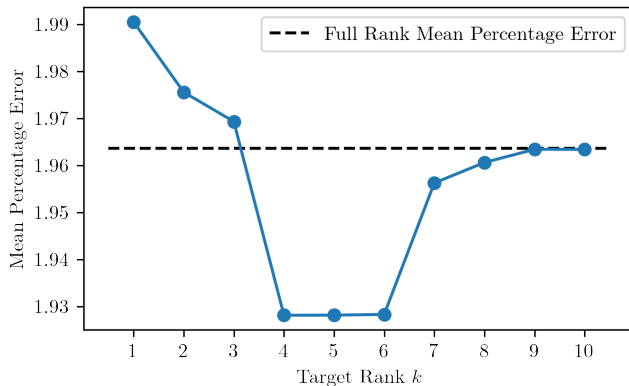


Table of Contents

Fast Multipole Methods (FMMs)

PyExaFMM

Research Context

Modern Architectures

Randomised SVD Compression

Modern Architectures

- Program for shared and distributed memory paradigms
[5, 4, 6]
- Take advantage of heterogenous architectures (GPU/CPU)
[4, 6]

Randomised SVD Compression

- Reduce cost of wastefully computing full-rank SVD, only to discard all but first k singular values/vectors
- Can be programmed with shared memory paradigm [1, 5]

References I

- [1] Apache mahout - distributed linear algebra.
- [2] L Greengard and V Rokhlin.
A fast algorithm for particle simulations.
Journal of Computational Physics, 73(2):325 – 348, 1987.
- [3] Denis Zorin Lexing Ying, George Biros.
A kernel-independent adaptive fast multipole algorithm in two and three dimensions.
Journal of Computational Physics, 196(2):591–626, 2004.
- [4] Dhairya Malhotra and George Biros.
Pvfm: A parallel kernel independent fmm for particle and volume potentials.
Communications in Computational Physics, 18(3):808–830, 2015.
- [5] Per-Gunnar Martinsson Nathan Halko and Joel Tropp.
Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions.
SIAM Review, 53:217–288, 01 2011.
- [6] Lorena A. Barba Rio Yokota.
Exafmm user's manual, 2011.